

Set Up Network and HTTP Load Balancers



Google Cloud Self-Paced Labs

Overview

In this hands-on lab, you'll learn the differences between a network load balancer and a HTTP load balancer, and how to set them up for your applications running on Google Compute Engine virtual machines.

There are several ways you can [load balance in Google Cloud Platform](#). This lab takes you through the setup of the following load balancers.:

- L3 [Network Load Balancer](#)
- L7 [HTTP\(s\) Load Balancer](#)

Students are encouraged to type the commands themselves, which helps in learning the core concepts. Many labs include a code block that contains the required commands. You can easily copy and paste the commands from the code block into the appropriate places during the lab.

What you'll do

- Setup a network load balancer.
- Setup a HTTP(s) load balancer.
- Get hands-on experience learning the differences between network load balancers and HTTP load balancers.

Prerequisites

Familiarity with standard Linux text editors such as `vim`, `emacs`, or `nano` is helpful.

What you need

To complete this lab, you need:

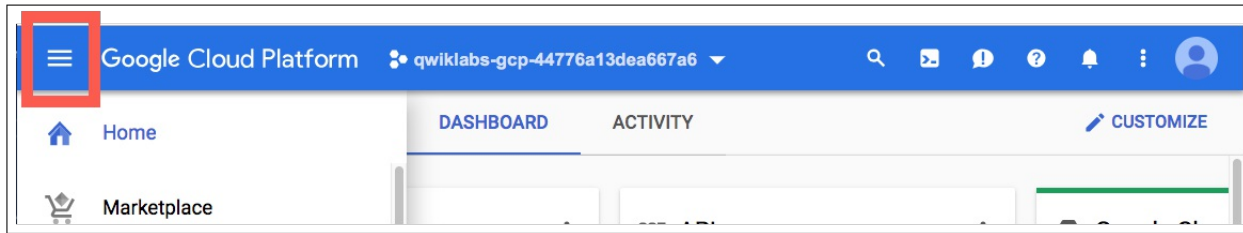
- Access to a standard internet browser (Chrome browser recommended).
- Time to complete the lab.

How to start your lab and sign in to the Console

- Open <https://console.cloud.google.com/>
- Enter login credentials

After a few moments, the GCP console opens in this tab.

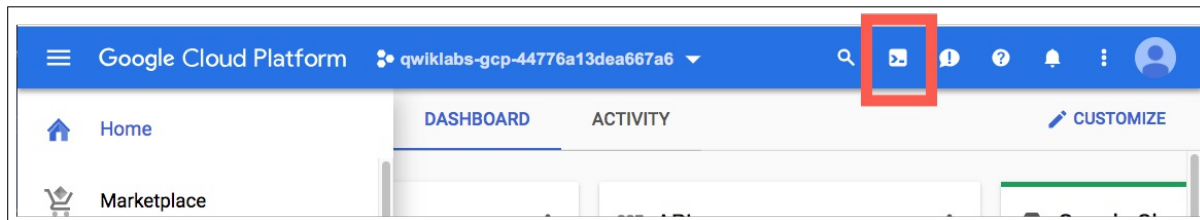
Note: You can view the menu with a list of GCP Products and Services by clicking the **Navigation menu** at the top-left, next to "Google Cloud Platform".



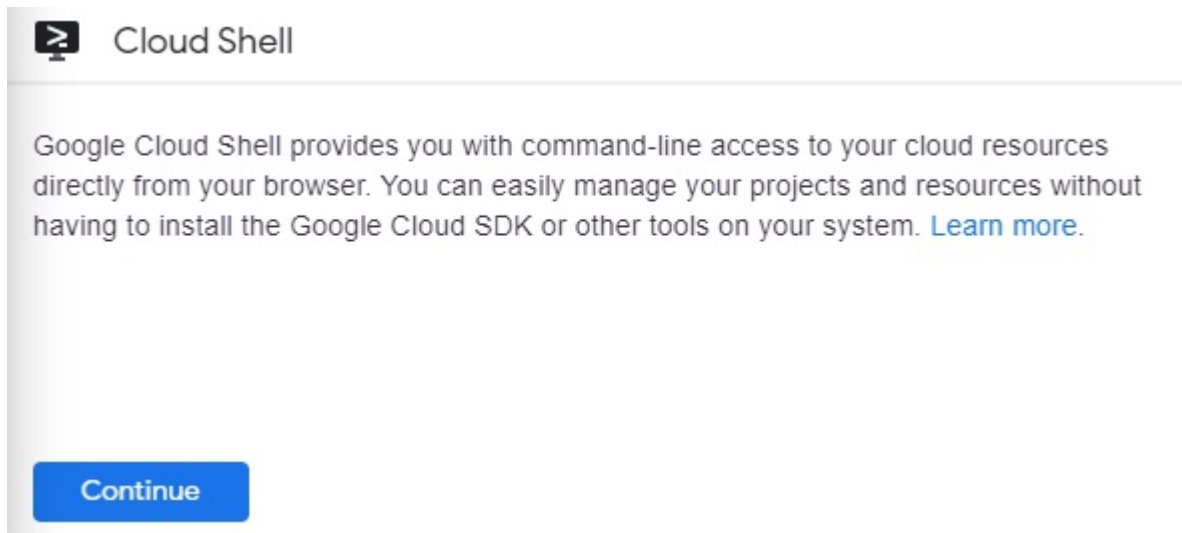
Activate Google Cloud Shell

Google Cloud Shell is a virtual machine that is loaded with development tools. It offers a persistent 5GB home directory and runs on the Google Cloud. Google Cloud Shell provides command-line access to your GCP resources.

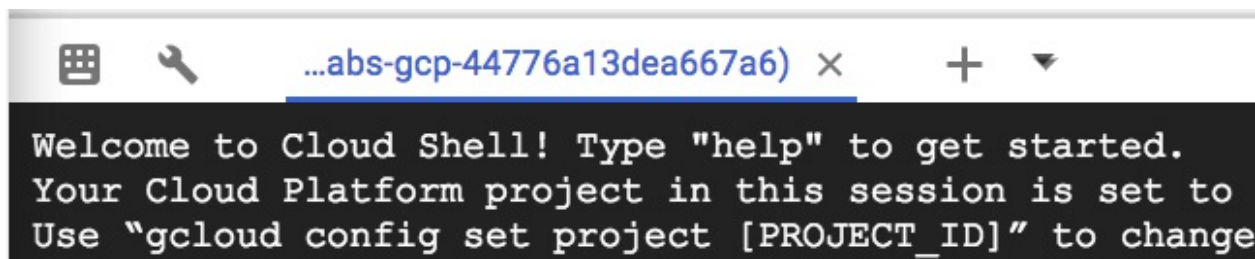
1. In GCP console, on the top right toolbar, click the Open Cloud Shell button.



2. Click **Continue**.



It takes a few moments to provision and connect to the environment. When you are connected, you are already authenticated, and the project is set to your `PROJECT_ID`. For example:



gcloud is the command-line tool for Google Cloud Platform. It comes pre-installed on Cloud Shell and supports tab-completion.

You can list the active account name with this command:

```
gcloud auth list
```

Output:

```
Credentialed accounts:
- <myaccount>@<mydomain>.com (active)
```

Example output:

```
Credentialed accounts:
- google1623327_student@testlabs.net
```

You can list the project ID with this command:

```
gcloud config list project
```

Output:

```
[core]
project = <project_ID>
```

Example output:

```
[core]
project = testlabs-gcp-44776a13dea667a6
```

Full documentation of **gcloud** is available on [Google Cloud gcloud Overview](#).

Set the default region and zone for all resources

In Cloud Shell, set the default zone:

```
gcloud config set compute/zone us-central1-a
```

Set the default region:

```
gcloud config set compute/region us-central1
```

Learn more about choosing zones and regions here: [Regions & Zones documentation](#).

Note: When you run `gcloud` on your own machine, the `config` settings persist across sessions. In Cloud Shell you need to set this for every new session or reconnection.

Create multiple web server instances

To simulate serving from a cluster of machines, create a simple cluster of Nginx web servers to serve static content using [Instance Templates](#) and [Managed Instance Groups](#).

Instance Templates define the look of every virtual machine in the cluster (disk, CPUs, memory, etc). Managed Instance Groups instantiate a number of virtual machine instances using the Instance Template.

To create the Nginx web server clusters, create the following:

- A startup script to be used by every virtual machine instance to setup Nginx server upon startup
- An instance template to use the startup script
- A target pool
- A managed instance group using the instance template

Still in Cloud Shell, create a startup script to be used by every virtual machine instance. This script sets up the Nginx server upon startup:

```
cat << EOF > startup.sh
#!/bin/bash
apt-get update
apt-get install -y nginx
service nginx start
sed -i -- 's/nginx/Google Cloud Platform - '"$HOSTNAME"'/' /var/www/html/index.nginx-debian.html
EOF
```

Create an instance template, which uses the startup script:

```
gcloud compute instance-templates create nginx-template \
  --metadata-from-file startup-script=startup.sh
```

(Output)

```
Created [...].
NAME                MACHINE_TYPE  PREEMPTIBLE  CREATION_TIMESTAMP
nginx-template      n1-standard-1                2015-11-09T08:44:59.007-08:00
```

Create a target pool. A target pool allows a single access point to all

the instances in a group and is necessary for load balancing in the future steps.

```
gcloud compute target-pools create nginx-pool
```

(Output)

```
Created [...].
NAME          REGION    SESSION_AFFINITY  BACKUP  HEALTH_CHECKS
nginx-pool    us-central1
```

Create a managed instance group using the instance template:

```
gcloud compute instance-groups managed create nginx-group \
  --base-instance-name nginx \
  --size 2 \
  --template nginx-template \
  --target-pool nginx-pool
```

(Output)

```
Created [...].
NAME          LOCATION    SCOPE  BASE_INSTANCE_NAME  SIZE  TARGET_SIZE  INSTANCE_TEMPLATE  AUTOSCALED
nginx-group    us-central1-a  zone   nginx               0     2            nginx-template     no
```

This creates 2 virtual machine instances with names that are prefixed with `nginx-`. This may take a couple of minutes.

List the compute engine instances and you should see all of the instances created:

```
gcloud compute instances list
```

(Output)

NAME	ZONE	MACHINE_TYPE	PREEMPTIBLE	INTERNAL_IP	EXTERNAL_IP	STATUS
nginx-7wvi	us-central1-a	n1-standard-1		10.240.X.X	X.X.X.X	RUNNING
nginx-9mwd	us-central1-a	n1-standard-1		10.240.X.X	X.X.X.X	RUNNING

Now configure a firewall so that you can connect to the machines on port 80 via the `EXTERNAL_IP` addresses:

```
gcloud compute firewall-rules create www-firewall --allow tcp:80
```

You should be able to connect to each of the instances via their external IP addresses via `http://EXTERNAL_IP/` shown as the result of running the previous command.

Check your lab progress. Click **Check my progress** below to verify that you've created a group of web servers.

Create a group of web servers

Create a Network Load Balancer

Network load balancing allows you to balance the load of your systems based on incoming IP protocol data, such as address, port, and protocol type. You also get some options that are not available with HTTP(S) load balancing. For example, you can load balance additional TCP/UDP-based protocols such as SMTP traffic. And if your application is interested in TCP-connection-related characteristics, network load balancing allows your app to inspect the packets, where HTTP(S) load balancing does not.

For more information, see [Setting Up Network Load Balancing](#).

Create an L3 network load balancer targeting your instance group:

```
gcloud compute forwarding-rules create nginx-lb \
  --region us-central1 \
  --ports=80 \
  --target-pool nginx-pool
```

(Output)

```
Created [https://www.googleapis.com/compute/v1/projects/...].
```

List all Google Compute Engine forwarding rules in your project.

```
gcloud compute forwarding-rules list
```

(Output)

NAME	REGION	IP_ADDRESS	IP_PROTOCOL	TARGET
nginx-lb	us-central1	X.X.X.X	TCP	us-central1/targetPools/nginx-pool

You can then visit the load balancer from the browser

`http://IP_ADDRESS/` where `IP_ADDRESS` is the address shown as the result of running the previous command.

Check your lab progress. Click **Check my progress** below to verify that you've created an L3 Network Load Balancer that points to the web servers.

Create an L3 Network Load Balancer that points to the web servers

Create a HTTP(s) Load Balancer

HTTP(S) load balancing provides global load balancing for HTTP(S) requests destined for your instances. You can configure URL rules that route some URLs to one set of instances and route other URLs to other instances. Requests are always routed to the instance group that is closest to the user, provided that group has enough capacity and is appropriate for the request. If the closest group does not have enough capacity, the request is sent to the closest group that does have capacity.

Learn more about the [HTTP\(s\) Load Balancer](#) in the documentation.

First, create a [health check](#).

Health checks verify that the instance is responding to HTTP or HTTPS traffic:

```
gcloud compute http-health-checks create http-basic-check
```

(Output)

```
Created [https://www.googleapis.com/compute/v1/projects/...].
NAME          HOST PORT REQUEST_PATH
http-basic-check 80      /
```

Define an HTTP service and map a port name to the relevant port for the instance group. Now the load balancing service can forward traffic to the named port:

```
gcloud compute instance-groups managed \
  set-named-ports nginx-group \
  --named-ports http:80
```

(Output)

```
Updated [https://www.googleapis.com/compute/v1/projects/...].
```

Create a [backend service](#):

```
gcloud compute backend-services create nginx-backend \
  --protocol HTTP --http-health-checks http-basic-check --global
```

(Output)

```
Created [https://www.googleapis.com/compute/v1/projects/...].
NAME          BACKENDS PROTOCOL
nginx-backend      HTTP
```

Add the instance group into the backend service:

```
gcloud compute backend-services add-backend nginx-backend \
```

```
--instance-group nginx-group \  
--instance-group-zone us-central1-a \  
--global
```

(Output)

```
Updated [https://www.googleapis.com/compute/v1/projects/...].
```

Create a default URL map that directs all incoming requests to all your instances:

```
gcloud compute url-maps create web-map \  
--default-service nginx-backend
```

(Output)

```
Created [https://www.googleapis.com/compute/v1/projects/...].  
NAME      DEFAULT_SERVICE  
Web-map   nginx-backend
```

To direct traffic to different instances based on the URL being requested, see [content-based routing](#).

Create a target HTTP proxy to route requests to your URL map:

```
gcloud compute target-http-proxies create http-lb-proxy \  
--url-map web-map
```

(Output)

```
Created [https://www.googleapis.com/compute/v1/projects/...].  
NAME      URL_MAP  
http-lb-proxy web-map
```

Create a global forwarding rule to handle and route incoming requests. A forwarding rule sends traffic to a specific target HTTP or HTTPS proxy depending on the IP address, IP protocol, and port specified. The global forwarding rule does not support multiple ports.

```
gcloud compute forwarding-rules create http-content-rule \  
--global \  
--target-http-proxy http-lb-proxy \  
--ports 80
```

(Output)

```
Created [https://www.googleapis.com/compute/v1/projects/...].
```


After creating the global forwarding rule, it can take several minutes for your configuration to propagate.

```
gcloud compute forwarding-rules list
```

(Output)

NAME	REGION	IP_ADDRESS	IP_PROTOCOL	TARGET
http-content-rule		X.X.X.X	TCP	http-lb-proxy
nginx-lb	us-central1	X.X.X.X	TCP	us-central1/...

Take note of the http-content-rule IP_ADDRESS for the forwarding rule.

From the browser, you should be able to connect to `http://IP_ADDRESS/`. It may take three to five minutes. If you do not connect, wait a minute then reload the browser.

Check your lab progress. Click **Check my progress** below to verify that you've created an L7 HTTP(S) Load Balancer.

Create an L7 HTTP(S) Load Balancer

Test your knowledge

Test your knowledge about Google cloud Platform by taking our quiz. (Please select multiple correct options if necessary.)

Congratulations!

You built a network load balancer and a HTTP(s) load balancer. You practiced with Instance templates and Managed Instance Groups. You are well on your way to having your Google Cloud Platform project monitored with Cloud Monitoring.

Next Steps / Learn More

- Deploy more resources to your project, and see them get monitored
- Add your shiny new GCP Essentials badge to your resume!